



TASK

Databases and the MongoDB Database

Visit our website

Introduction

WELCOME TO THE DATABASES AND THE MONGODB DATABASE TASK!

Dynamic web applications rely on data. Storing data about your users is a good starting point if you want to make a web application more dynamic. In this task, you will learn about various types of databases. You will also be introduced to MongoDB (a very popular database), and shown how to use MongoDB to create a database-as-a-service architecture or implementation.


DATABASES

Data is core to software development. Programs are designed to manipulate, create and visualise data. Therefore, it is important for all software developers and web developers to be able to access, manipulate, and store data.

Databases are used to store data. A database is simply a large container of data with the ability to order the data in multiple ways while providing easy access to the data itself. Web developers often have to manipulate the data stored in databases. For example, you may need to store your users' usernames, passwords, names, addresses, telephone numbers, etc. Therefore, full-stack web developers need to be able to work proficiently with databases.

DATA VS INFORMATION

In order to properly understand databases, you must first understand the difference between data and information. Simply put, data are **raw** facts. The word raw indicates that the facts have not yet been processed to create meaning. Information, on the other hand, is the result of processing raw data to make it meaningful. Data processing can be as simple as organising data to reveal patterns, or as complex as making forecasts or drawing inferences using statistical modelling.



The production of accurate, relevant, and timely information is the key to good decision-making. In turn, good decision-making is the key to a business's survival in a competitive global environment. Timely and useful information requires accurate data, which must be captured properly and stored in a format that is easy to access and process. The data environment should be carefully managed.

DATABASE MANAGEMENT SYSTEM

A database can be thought of as a well-organised electronic filing cabinet where powerful software, known as a database management system (DBMS), helps manage the contents of the cabinet. A database management system is a collection of programs that manage the database structure and control access to the data stored in the database.

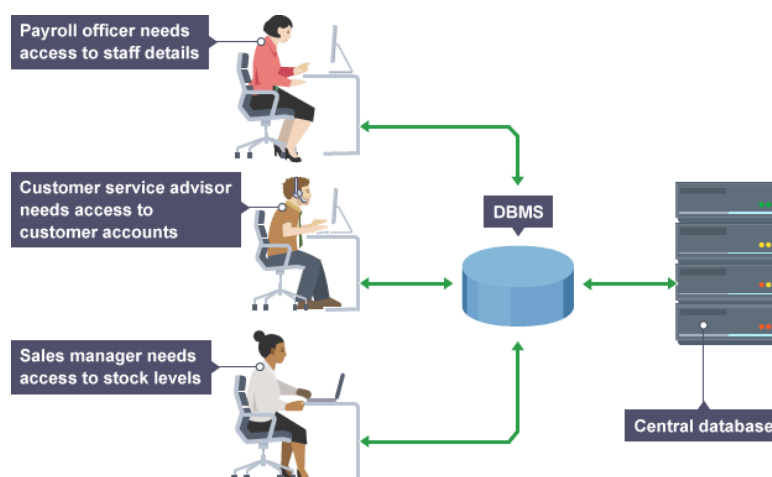


Image source:

<https://www.chrisbell.com/SNHU/IT-510-Advanced-Information-Technology/database-management-system-DBMS-relationships.php>

The illustration above shows how the DBMS serves as an intermediary between the user and the database. The DBMS receives all application requests and translates them into the complex operations required to fulfil those requests.

Much of the database's internal complexity is hidden from the application programs and end users by the DBMS. There are some very important advantages to having a DBMS between the end user's application and the database. Firstly, the DBMS allows the data in the database to be shared among multiple applications or users. Secondly, the DBMS integrates many different users' views of the data into a single data repository.

The DBMS helps make data management much more efficient and effective and provides advantages such as:

- **Improved data sharing:** the DBMS helps create an environment in which end users have better access to more and better-managed data.
- **Better data integration:** an integrated view of the organisation's operations and a clearer view of the big picture is promoted by wider access to well-managed data.
- **Minimised data inconsistency:** data inconsistency occurs when different versions of the same data appear in different places. A properly designed database greatly reduces the probability of data inconsistency.
- **Improved data access:** a query is a specific request for data manipulation (e.g. to read or update the data) sent to the DBMS. The DBMS makes it possible to produce quick answers to spur-of-the-moment queries.
- **Improved decision-making:** better-quality information (on which to base decisions) is generated due to better-managed data and improved data access.
- **Increased end-user productivity:** the availability of data and the tools that transform data into usable information encourages end users to make quick, informed decisions.

TYPES OF DATABASES

There are many different types of databases. These databases can be classified according to the number of users supported, where the data are located, the type of data stored, the intended data usage and the degree to which the data are structured.

A database can be classified as either **single-user** or **multi-user**. A single-user database only supports one user at a time. This means that if user A is using the

database, users B and C must wait until user A is done. A desktop database is a single-user database that runs on a personal computer. A multi-user database, on the other hand, supports multiple users at the same time. A workgroup database is a multi-user database that supports a relatively small number of users (usually less than 50) or a specific department within an organisation. When a multi-user database supports many users (more than 50) and is used by the entire organisation across many departments, it is known as an enterprise database.

Databases can also be classified based on location. A **centralised database** supports data located at a single site, whereas a **distributed database** supports data distributed across several different sites.

Another popular way of classifying databases is based on how they will be used and on the time sensitivity of the information gathered from them. An example of this is an **operational database**, which is a database that is designed to primarily support a company's day-to-day operations. Operational databases are also known as online transaction processing (OLTP), transactional, or production databases.

The degree to which data is structured is another way of classifying databases. Data that exist in their original, or raw, state are known as **unstructured data**. In other words, they are in the format in which they were collected. **Structured data** are the result of formatting unstructured data to facilitate the storage, use, and generation of information. You apply structure based on the type of processing that you intend to perform on the data. For example, imagine that you have a stack of printed invoices. If you just want to store these invoices so that you are able to retrieve them or display them later, you can scan them and save them in a graphical format. However, if you want to derive information from them, such as monthly totals or average sales, having the invoices in a graphical format will not be useful. You could instead store the invoice data in a structured spreadsheet format so that you can perform the desired computations.

Analytical databases focus on storing historical data and business metrics used exclusively for tactical or strategic decision-making. They typically comprise two components: a data warehouse, and an online analytical processing (OLAP) front end. Analytical databases allow the end user to perform advanced data analysis of business data using sophisticated tools. By contrast, a data warehouse focuses on storing data used to generate the information required to make tactical or strategic decisions.

There are many ways of storing data, each with its own benefits and limitations. Therefore, databases can be designed differently to meet different needs. We are

going to consider two of the most popular types of databases: relational databases and NoSQL databases.

Relational databases

Relational databases store information about different entities and the relationships between them. The image below is an example of an entity-relationship diagram (ERD) that is used to describe the relationships between certain entities.

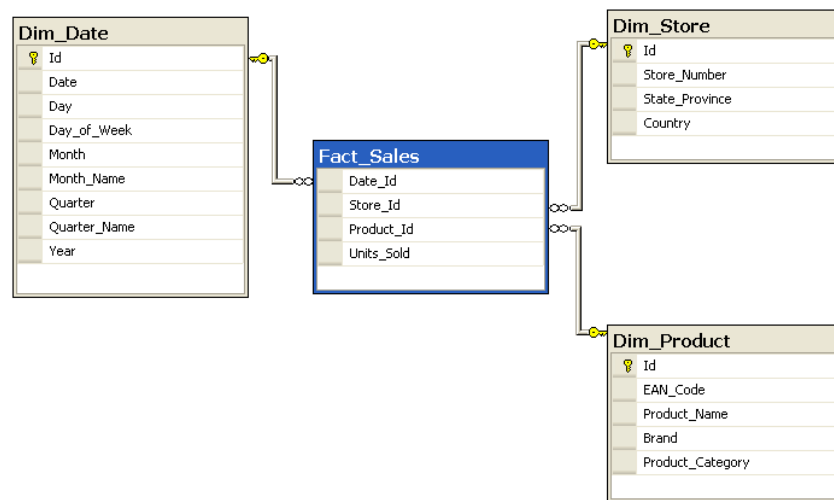


Image source:

https://en.wikipedia.org/wiki/Star_schema#/media/File:Star-schema-example.png

In a relational database, an entity is a table that stores all the data about a certain thing. For example, you may want to store information about all your customers in a database. You would then create a customer table (customer entity), which could look something like the one shown below:

C_NAME	C_PHONE	C_ADDRESS	C_POST CODE	A_NAME	A_PHONE	TP	AMT	REN
Alfred Smith	082 345 2341	207 Willow St, Port Elizabeth	6390	Leah Hahn	084 259 2073	T1	R100.00	05-Apr-2021
Kathy Dunne	083 567 9012	556 Bad St, Cape Town	7100	Alex Alby	085 785 3938	S2	R250.00	16-Jun-2021
Paul Farris	076 782 1232	2148 High St, Benoni	1522	Leah Hahn	084 259 2073	T2	R850.00	22-Sep-2021

Customer entity fields:

C_NAME = customer name

C_PHONE = customer phone

C_ADDRESS = customer address

C_POSTCODE = customer postcode

A_NAME = agent name

A_PHONE = agent phone

TP = insurance type

AMT = insurance policy amount in thousands of R

REN = Insurance renewal date

The **CUSTOMER** table contains three records. Each record is composed of nine fields: **C_NAME**, **C_PHONE**, **C_ADDRESS**, **C_POSTCODE**, **A_NAME**, **A_PHONE**, **TP**, **AMT**, and **REN**. Each record describes a specific customer; each customer is an instance of the customer entity.

If you were designing a database for a store, you might also have a product entity that stores all the information about the products you sell. Your database would then contain a product table and a customer table. The database would also save information about the relationship between the two entities. For example, it would store information about which products a particular customer bought on a particular date. However, how this is implemented in a relational database is beyond the scope of this course.

NoSQL databases

A problem with relational databases is that their performance degrades as the volume of data increases. Many web applications have to store massive amounts of data. Imagine the volume of data that companies like Amazon and Google have to store, for example. Addressing this problem led to the development of NoSQL databases. You are using a NoSQL database every time you search for a product on Amazon, watch a video on Youtube, or send a message to a friend on Facebook. NoSQL databases generally have the following characteristics:

- They are not based on the relational model.
- They support distributed database architectures, i.e. servers in different areas.
- They provide high scalability, high availability, and fault tolerance.
- They support very large amounts of sparse data (data in which there are a lot of attributes but many of these are not populated).
- They are geared toward performance rather than transaction consistency.

There are several types of NoSQL databases. Some of these are briefly described below:

1. **Key-value store databases:** This is the simplest form of NoSQL database. Every item in the database is stored as a key (used to identify the value) and its value.



Image source:

<https://www.slideshare.net/arangodb/introduction-to-column-oriented-databases>

2. **Column-oriented databases:** Like key-value store databases, a key is used to identify values, but instead of the key identifying only one value, it can be used to identify multiple values.




Image source:

<https://www.slideshare.net/arangodb/introduction-to-column-oriented-databases>

3. **Document store databases:** With this type of database, a key is used to identify a particular document. Documents are stored in recognised formats like XML, JSON, PDF, etc.





Image source:

<https://www.slideshare.net/arangodb/introduction-to-column-oriented-databases>

4. **Graph databases:** This database uses the graph data structure to store data. Graphs contain nodes (A, B, C, and D in the image below) and edges (the lines connecting the nodes, so AB, BD, BC, and DC in the image below). In graph databases, nodes are objects and edges are the relationships between these objects. Social networking applications, like Facebook, often store data using graph databases because this model is good for tracking the relationships between objects.



5. **Object-oriented databases:** These databases combine object-oriented programming (OOP) and database principles. These databases are tied to specific programming languages.

MONGODB

In this bootcamp you will be working with MongoDB, a document store and NoSQL database. MongoDB is made up of collections and documents.

- **Collection:** A collection is a group of documents. It is similar to an entity or table when working with relational databases.
- **Documents:** In relational databases, records are stored in tables. An example of a record in the CUSTOMER table we considered earlier would be Alfred Smith. MongoDB uses documents instead of records (or rows in a table) to store data (i.e. with a MongoDB database, Alfred Smith's data would be stored in a document instead of in a row in the CUSTOMER table). MongoDB uses BSON documents. BSON stands for Binary JSON. BSON uses JSON files and stores type information. JSON files are just text information; this means that it has to be parsed if you want to program with it. As BSON stores type information, it is quicker and more efficient to use than JSON (see an example below).

```

{
  name : "Joe Drumgoole",
  title : "Director of Developer Advocacy",
  Address : {
    address1 : "Latin Hall",
    address2 : "Golden Lane",
    eircode : "D09 N623",
  },
  expertise: [ "MongoDB", "Python", "Javascript" ],
  employee_number : 320,
  location : [ 53.34, -6.26 ]
}

```

Strings

Nested Document

Array

Integer

Geo-spatial Coordinates




Extra resource

Watch [this video](#) in which MongoDB's CTO and Co-Founder, Eliot Horowitz, gives a comprehensive introduction to MongoDB in under five minutes.

MONGODB IN A FULL-STACK WEB APPLICATION

Before we start using MongoDB, it is important to first see how it fits into full-stack web development and what other tools are needed to get it to work properly.



Consider the image above. As you already know, to use your app, clients will interact with a web server that will be running Node.js. Node.js will route all requests and perform whatever server-side logic is needed by our app. If our user wants to access, add, or change information that needs to persist, they will need access to the MongoDB database. An important component of MongoDB is

Mongo, its administrative shell. Mongo is a program that allows you to execute instructions on the database from a command line interface. Mongo allows you to use the MongoDB query language.

For Node.js to be able to communicate with MongoDB, it also makes use of *MongoDB drivers*. The official MongoDB driver can be installed using NPM (more on this later).

MONGODB AS A SERVICE

A few years ago, in order to run a database, you had to have a server (or servers) with all the necessary software installed and configured. You would also need someone who would act as the database administrator. In many organisations, this is still the case and this may be a justified expense and effort. Today though, there is an attractive alternative: database as a service.

Since cloud computing has become more popular, there are more cloud-based options for developers. You can host your web app on the cloud (via Heroku, Azure, etc.) instead of on your own dedicated server. You can also use a database hosted by a cloud service provider (as in the case of MongoDB's Atlas) rather than going through the pain of setting up and maintaining your own database server. Below are some key benefits of this approach:

1. It is often cheaper than having your own database server because you only pay for what you use.
2. The cloud service provider deals with all the hassle of managing the configuration, backup, maintenance, and security of the database server.
3. It is quick and easy to start using a database with minimal configuration.

In this bootcamp, we will be using MongoDB's database as a service solution: Atlas.

In this task, you will create your first database using MongoDB. Before you can do this though, you are first going to:

1. download and install MongoDB on your local machine so that you can use Mongo, the administrative shell,
2. use Atlas to create and host a MongoDB on the cloud, and
3. use Mongo to access and manipulate your database cluster on Atlas.

INSTALL MONGODB

We will be installing MongoDB's free Community Server for this course. An alternative is the Enterprise Server which you can download as a free trial. Follow the steps linked below to get started with MongoDB.

- **[MongoDB Community Edition Installation Steps](#)**

The steps to set up and configure MongoDB differ slightly depending on the OS you are using. The installation instructions given above contain instructions for running MongoDB on your specific operating system.

Your Mongo shell is now ready to be used to connect to a database server. MongoDB Atlas will provide the platform and infrastructure we need for a database server on the cloud.

SETUP MONGODB ATLAS

As previously stated, in this bootcamp, we will be using MongoDB's Database as a service solution: Atlas. To get a quick (two-minute) overview of what Atlas is and why we are using it, watch this **[short video](#)**.

To configure MongoDB Atlas, do the following:

- Go **[here](#)** and enter your information to get started with Atlas.
- You will then be taken to the "Database Deployments" page:
 - Click on the "Build a Database" button on screen.
 - You'll be redirected to a new screen.
 - Under 'Cloud provider & Region' select AWS and any free tier region.
 - Under 'Cluster Tier' select the free M0 option.
 - You can rename your cluster under 'Cluster Name'.



Deploy your database

Use a template below or set up [advanced configuration options](#). You can also edit these configuration options once the cluster is created.

M10 \$0.08/hour For production applications with sophisticated workload requirements. <table><tr><td>STORAGE</td><td>RAM</td><td>vCPU</td></tr><tr><td>10 GB</td><td>2 GB</td><td>2 vCPUs</td></tr></table>	STORAGE	RAM	vCPU	10 GB	2 GB	2 vCPUs	SERVERLESS \$0.10/1M reads For application development and testing, or workloads with variable traffic. <table><tr><td>STORAGE</td><td>RAM</td><td>vCPU</td></tr><tr><td>Up to 1 TB</td><td>Auto-scale</td><td>Auto-scale</td></tr></table>	STORAGE	RAM	vCPU	Up to 1 TB	Auto-scale	Auto-scale	M0 FREE For learning and exploring MongoDB in a cloud environment. <table><tr><td>STORAGE</td><td>RAM</td><td>vCPU</td></tr><tr><td>512 MB</td><td>Shared</td><td>Shared</td></tr></table>	STORAGE	RAM	vCPU	512 MB	Shared	Shared
STORAGE	RAM	vCPU																		
10 GB	2 GB	2 vCPUs																		
STORAGE	RAM	vCPU																		
Up to 1 TB	Auto-scale	Auto-scale																		
STORAGE	RAM	vCPU																		
512 MB	Shared	Shared																		

Provider



Region

★ Recommended region ⓘ



Name

You cannot change the name once the cluster is created.

HyperionDev123456

FREE



Free forever! Your M0 cluster is ideal for experimenting in a limited sandbox. You can upgrade to a production cluster anytime.

[I'll deploy my database later](#)

[Access Advanced Configuration](#)

- Click on the 'Create' button at the bottom of the page to create your cluster.

Once you have created your cluster, you will be taken to a page similar to the one shown below. You can change the username and other settings for your database user as shown in the image.

Atlas HyperionDev Access Manager Billing

Project 0 Data Services App Services Charts

Overview

DEPLOYMENT

Database

Data Lake PREVIEW

SERVICES

Triggers

Data API

Data Federation

Search

SECURITY

Quickstart

Database Access

Network Access

Advanced

New On Atlas 4

Goto

HYPERIONDEV > PROJECT 0

Security Quickstart

To access data stored in Atlas, you'll need to create users and set up network security controls. [Learn more about security setup](#)

1 How would you like to authenticate your connection?

Your first user will have permission to read and write any data in your project.

Username and Password Certificate

We autogenerated a username and password for your first database user in this project using your MongoDB Cloud registration information.

Create a database user using a username and password. Users will be given the *read and write to any database privilege* by default. You can update these permissions and/or create additional users later. Ensure these credentials are different to your MongoDB Cloud username and password.

Username

HyperionDev_Student1

Password

3geen8FP0hJyhxE

Autogenerate Secure Password

Copy

Create User

Your cluster has finished provisioning. to connect from?

Enable access for any network(s) that need to read and write data to your cluster.

ADVANCED

You have successfully created a database cluster. You will now use the administrative shell, Mongo, to add a database to this cluster. Before you do this though, there are some security settings that you need to tweak to make sure that you can access the Atlas cluster from your computer.

Security settings

One of the ways that MongoDB ensures security is by only allowing certain machines (IP addresses) access to your cluster. You can Whitelist your IP by scrolling down on the "Quickstart" tab and clicking "Add my current IP Address".

After also adding a user, click “Finish and Close” at the bottom to continue.

Atlas HyperionDev Access Manager Billing

Project 0 Data Services App Services Charts

Overview

DEPLOYMENT

Database

Data Lake PREVIEW

SERVICES

Triggers

Data API

Data Federation

Search

SECURITY

Quickstart

Database Access

Network Access

Advanced

New On Atlas 4

Goto

..... Autogenerate Secure Password Copy

Create User

2 Where would you like to connect from?

Enable access for any network(s) that need to read and write data to your cluster.

My Local Environment
Use this to add network IP addresses to the IP Access List. This can be modified at any time.

Cloud Environment ADVANCED
Use this to configure network access between Atlas and your cloud or on-premise environment. Specifically, set up IP Access Lists, Network Peering, and Private Endpoints.

Add entries to your IP Access List

Only an IP address you add to your Access List will be able to connect to your project's clusters.

IP Address	Description
<input type="text" value="Enter IP Address"/>	<input type="text" value="Enter description"/>

Add My Current IP Address

Add Entry

Finish and Close

System Status: All Good

©2023 MongoDB, Inc. Status Terms Privacy Atlas Blog Contact Sales

The IPs that are allowed to access your cluster are listed under the “Network Access” tab in the left hand navbar. If you currently have a dynamic IP address, you’ll need to allow access from anywhere.

To do so, click on the “Network Access” tab and then click on the “ADD IP ADDRESS” button. The following pop-up will then appear.

Add IP Access List Entry

Atlas only allows client connections to a cluster from entries in the project's IP Access List. Each entry should either be a single IP address or a CIDR-notated range of addresses. [Learn more](#).

ALLOW ACCESS FROM ANYWHERE

Access List Entry:

Comment:

☐ This entry is temporary and will be deleted in 6 hours Cancel Confirm

Click "ALLOW ACCESS FROM ANYWHERE" and click "Confirm".

The following entry should appear in your Network Access tab confirming you've allowed access from anywhere:

HYPERIONDEV > PROJECT 0

Network Access

IP Access List Peering Private Endpoint + ADD IP ADDRESS

You will only be able to connect to your cluster from the following list of IP Addresses:

IP Address	Comment	Status	Actions
0.0.0.0/0 (includes your current IP address)		Active	EDIT DELETE

It is not good practice to allow all IP addresses to access your database for obvious security reasons, but for the purposes of the next few tasks, we are going to ask you to do this. The reason for this is that you will need to give a mentor access to your database, and you won't know their IP address ahead of time. In professional practice, however, it is advisable to have a limited IP whitelist.



Take note:

You should remember to configure the IP whitelist when you deploy your app. Remember that you will ultimately write code to access your database in your Express app (which you will soon learn to do). Therefore, your application server will be making requests to your database. If you deploy your back-end app to the cloud (e.g. deploy to AWS) you do not necessarily know what the IP

address of your web server will be! This could obviously be a problem – if your web server IP address isn't added to the Atlas IP whitelist, your Express app won't be allowed to communicate with your database!

To address this problem you can use Render's Static Outbound IP Addresses. According to an [article by Render](#), "You can use these addresses to connect to IP-restricted environments outside your Render network". Outbound requests from those apps originate from a set of static IP addresses, which allows you to securely communicate with IP white-listed services on-premise (hosted on-site) or on other networks.

If your machine is protected by a firewall, you also have to ensure that this doesn't block access to Atlas. Atlas servers run on **port 27017** on Amazon AWS. Check [here](#) to see if this port is blocked on your machine or not. If [this page](#) doesn't load, your firewall is probably blocking port 27017. If it is blocked, make sure to unblock it before you proceed. How this is done will depend on the firewall you are using. Google the appropriate instruction to unblock port 27017 for the firewall you are running.

From the Security tab, you can also add users and manage the rights of the users you allow to access your database.

Manage users and teams

As the database administrator, you have to manage who is able to access your database and what they can do with your database. For your next MongoDB tasks, you are required to give a mentor access to your database. To do this, select "Access Manager" and then click on "Project Access" along with your current project displayed next to it on the drop-down. On the next screen, click on "Invite to Project". You can then invite a mentor to be a user of your database by entering an email address as indicated in the image below. Please use the email address: **allreviewers@hyperiondev.com**.

Atlas HyperionDev Access Manager Billing

Project 0 Data Services App Services Charts

PROJECT

- Settings
- Alerts
- Activity Feed
- Access Manager
- Support
- Integrations

HYPERIONDEV > PROJECT 0 > PROJECT ACCESS MANAGER

Invite to Project

Invite new or existing users, teams, and API keys to this project.

Search for existing users, teams, API key or invite new users via email address.

allreviewers@hyperiondev.com

Cancel Invite to Project

After pasting the email address, a suggestion should appear below it, and you'll need to click it in order to select the user.

Note: please ensure you give the mentor **full** access to your project.

Atlas HyperionDev Access Manager Billing

Project 0 Data Services App Services Charts

PROJECT

- Settings
- Alerts
- Activity Feed
- Access Manager
- Support
- Integrations

HYPERIONDEV > PROJECT 0 > PROJECT ACCESS MANAGER

Invite to Project

Invite new or existing users, teams, and API keys to this project.

Search for existing users, teams, API key or invite new users via email address.

Search by email address, team name, or public key

Give your members access permissions below.

allreviewers@hyperiondev.com

- Project Read Only, Project Owner, Project Search Index Editor, Project Data Access Read Only, Project Data Access Read/Write, Project Data Access Admin, Project Cluster Manager

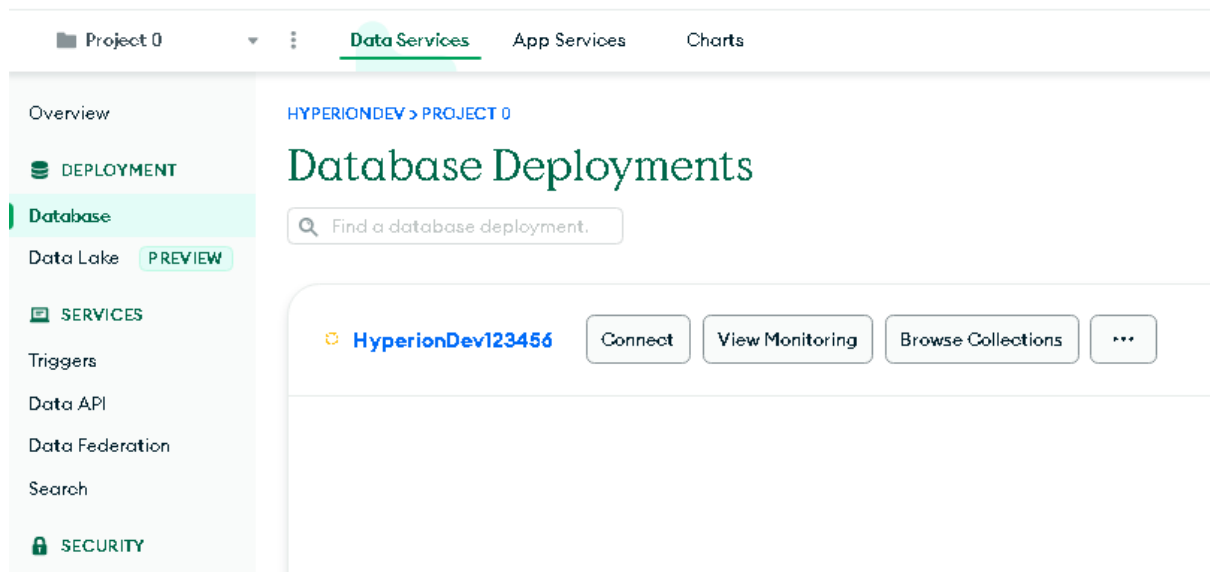
Remove

Cancel Invite to Project

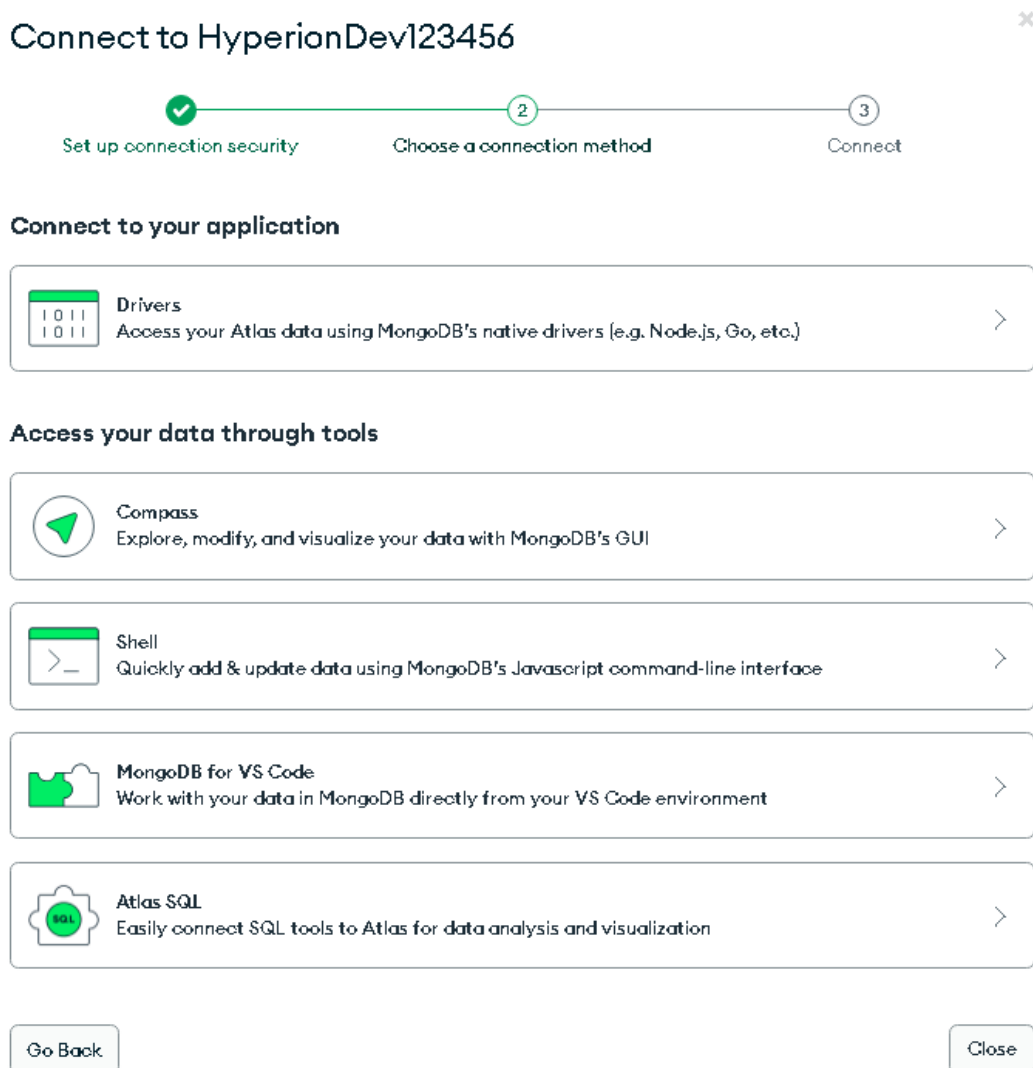
ACCESS THE DATABASE ON THE CLOUD USING THE MONGO SHELL

You are now ready to access the database server you have configured on MongoDB Atlas using the Mongo shell on your local machine. Remember that Mongo is the administrative shell used to run instructions on your MongoDB server.

To connect to the database server using Mongo, you need a connection string that specifies everything needed for this connection. Atlas provides this connection string for you. Select "Connect" as shown below, to find the connection string.



The following pop-up window will appear:



From the window above, select “Shell”. The window below should then appear. Select the version of Mongo Shell you have installed (for example in our case I selected: “mongosh (1.9 or later)”, but you may need to select the version that works for you).

Connect to HyperionDev123456 ✕



I don't have the MongoDB Shell installed

I have the MongoDB Shell installed

1. Select your mongo shell version

To check your Mongo shell version, run:
`mongosh --version` or `mongo --version`

mongosh (1.9 or later) ▼

2. Run your connection string in your command line

Use this connection string in your application

```
mongosh "mongodb+srv://hyperiondev123456.kobecqc.mongodb.net/" --apiVersion 1 --  
username HyperionDevSutend1
```

You will be prompted for the password for the Database User, **HyperionDevSutend1**. When entering your password, make sure all special characters are [URL encoded](#).

RESOURCES

[Add Data in the Shell](#)

[Troubleshoot Connections](#)

[Access your Database Users](#)

Go Back

Close

This will create a new connection string. Copy and paste the connection string that appears there into your command line interface. You should see output similar to that shown in the image below:

```
E:\>mongo "mongodb+srv://hyperion-f78fc.mongodb.net/test" --username hyperionDB
MongoDB shell version v3.6.2
Enter password:
connecting to: mongodb+srv://hyperion-f78fc.mongodb.net/test
2018-05-09T10:06:44.519+0200 I NETWORK [thread1] Starting new replica set monitor for hyperion-shard-00-00-f78fc.mongodb.net.:27017,hyperion-shard-00-01-f78fc.mongodb.net.:27017,hyperion-shard-00-02-f78fc.mongodb.net.:27017
2018-05-09T10:06:45.554+0200 I NETWORK [ReplicaSetMonitor-TaskExecutor-0] Successfully connected to hyperion-shard-00-00-f78fc.mongodb.net.:27017 (1 connections now open to hyperion-shard-00-00-f78fc.mongodb.net.:27017)
2018-05-09T10:06:45.568+0200 I NETWORK [thread1] Successfully connected to hyperion-shard-00-02-f78fc.mongodb.net.:27017 (1 connections now open to hyperion-shard-00-02-f78fc.mongodb.net.:27017 with a 5 second timeout)
2018-05-09T10:06:45.823+0200 I NETWORK [thread1] changing hosts to hyperion-shard-0/hyperion-shard-00-01-f78fc.mongodb.net.:27017,hyperion-shard-00-02-f78fc.mongodb.net.:27017,hyperion-shard-00-00-f78fc.mongodb.net.:27017,hyperion-shard-00-01-f78fc.mongodb.net.:27017,hyperion-shard-00-02-f78fc.mongodb.net.:27017
2018-05-09T10:06:46.838+0200 I NETWORK [ReplicaSetMonitor-TaskExecutor-0] Successfully connected to hyperion-shard-00-02-f78fc.mongodb.net.:27017 (1 connections now open to hyperion-shard-00-02-f78fc.mongodb.net.:27017 with a 5 second timeout)
2018-05-09T10:06:46.848+0200 I NETWORK [thread1] Successfully connected to hyperion-shard-00-00-f78fc.mongodb.net.:27017 (1 connections now open to hyperion-shard-00-00-f78fc.mongodb.net.:27017 with a 5 second timeout)
2018-05-09T10:06:48.132+0200 I NETWORK [ReplicaSetMonitor-TaskExecutor-0] Successfully connected to hyperion-shard-00-01-f78fc.mongodb.net.:27017 (1 connections now open to hyperion-shard-00-01-f78fc.mongodb.net.:27017 with a 5 second timeout)
MongoDB server version: 3.4.14
WARNING: shell and server versions do not match
MongoDB Enterprise hyperion-shard-0:PRIMARY>
```

You have connected to your MongoDB server hosted by Atlas! You are now able to use the Mongo shell to create and modify databases on your server.



Extra resource

Watch [this short video](#) to see how easily you can connect to your Atlas database to the Mongo shell.

CREATE A DATABASE

Once you can access your database server (run by Atlas), you can issue instructions using Mongo to change your database. We are going to create a database.

To do this, type the following using the Mongo shell: **use test** (where “test” is the name of the database). If the database does not already exist, this instruction will create it.

```
MongoDB Enterprise hyperion-shard-0:PRIMARY> use test
switched to db test
MongoDB Enterprise hyperion-shard-0:PRIMARY>
```

MONGODB COMPASS

You may have noticed that when you installed MongoDB, MongoDB Compass was also installed. Compass allows you to interface with your database. You should be able to connect to your database using Compass too. Give it a try.

QUIT MONGO

To quit mongo, type `quit()` into the Mongo shell.

Instructions

Take your time and work carefully through the Compulsory Task instructions below, one point at a time.

Compulsory Task

Follow these steps:

- Install MongoDB. See appropriate detailed installation instructions on MongoDB's download centre for details.
- Create a cluster on MongoDB Atlas.
- Add a mentor or code reviewer as a user to your cluster on Atlas, using the email address **allreviewers@hyperiondev.com**.
- Ensure that your firewall isn't blocking access to MongoDB.
- Connect to your cluster using the Mongo shell.
- Make a database called "test"
- Create a document called "myMongoDB" in which you include the following:
 - A screenshot that shows that you have added a mentor or code reviewer as a MongoDB user to your Atlas cluster.
 - A screenshot of your command line interface that shows how you have used the Mongo shell to connect to your MongoDB Atlas cluster.
 - A screenshot of your command line interface that shows that you have successfully created a database called "test" on your MongoDB Atlas cluster.

Completed the task(s)?

Ask an expert code reviewer to review your work!

[Review work](#)



Rate us

Share your thoughts

HyperionDev strives to provide internationally-excellent course content that helps you achieve your learning outcomes.

Think that the content of this task, or this course as a whole, can be improved, or think we've done a good job?

[Click here](#) to share your thoughts anonymously.

